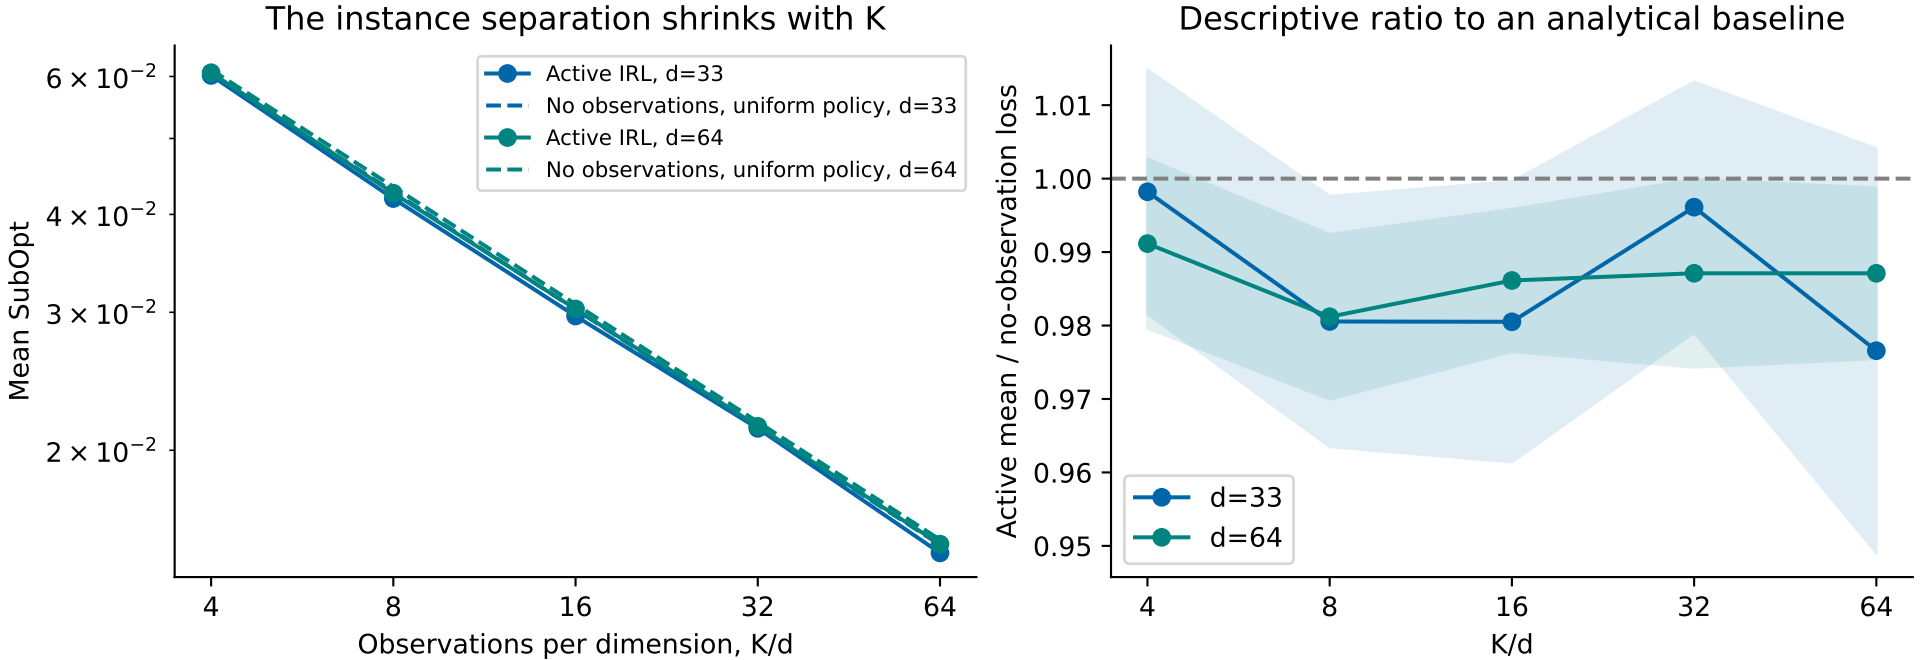
The instance separation shrinks with K
Mean SubOpt
6×10^{-2}
4×10^{-2}
3×10^{-2}
2×10^{-2}
Active IRL, d=33
No observations, uniform policy, d=33
Active IRL, d=64
No observations, uniform policy, d=64
4
8
16
32
64
Observations per dimension, K/d
Descriptive ratio to an analytical baseline
Active mean / no-observation loss
1.01
1.00
0.99
0.98
0.97
0.96
0.95
d=33
d=64
4
8
16
32
64
K/d
Shading: pointwise 95% intervals over 16 coordinate-balanced replicates.